

Designing High-Confidence Embedded Systems

Operating Systems and Real-Time

Gabe Palmer



Embedded Systems

Computers *embedded* in the physical world

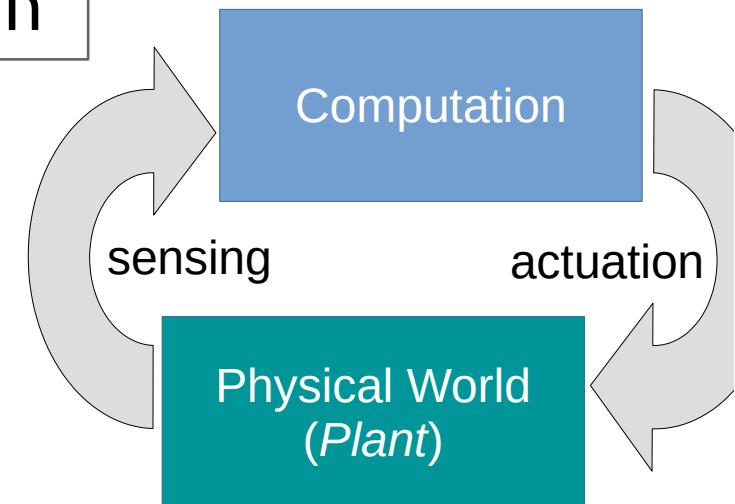
Computations **responsible** for physical interaction

Errors/Attacks → potential physical catastrophe



Embedded Systems

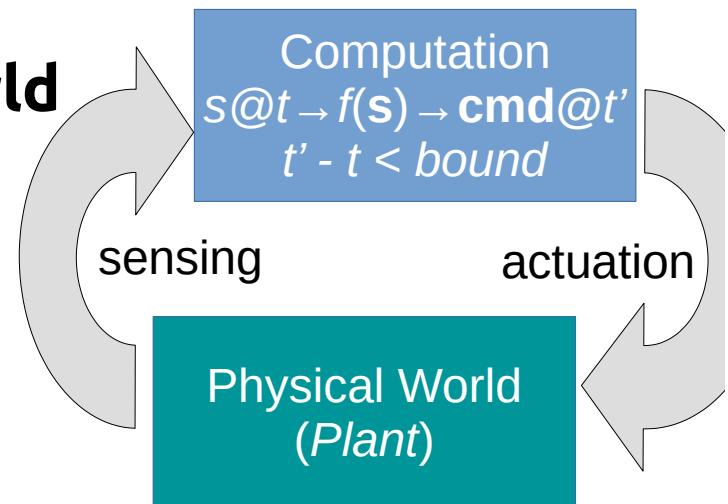
- Computers that interact with physical world
 - Sensors + actuators + computation
- Replace mechanical w/ digital
 - Microwaves, toys
 - Cars, airplanes



Real-Time/Cyber-Physical Systems

Embedded systems + **deadlines**

- The computation must adhere to the **timing constraints of the physical world**
- Send command to actuator within a **bounded time** from reading sensor



Real-Time/Cyber-Physical Systems

Embedded systems + **deadlines**

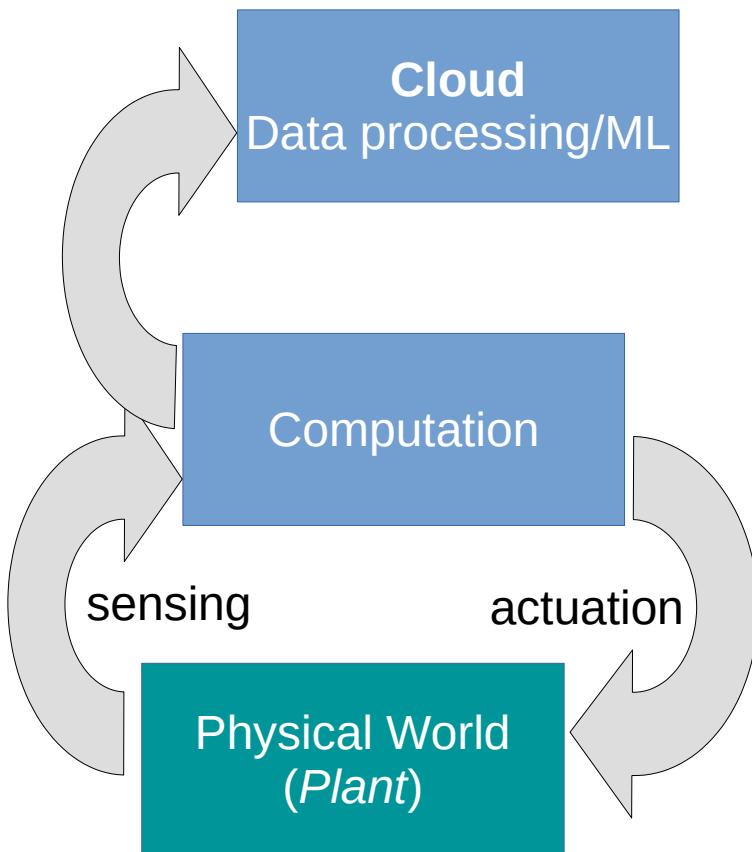
- The computation must adhere to the **timing constraints of the physical world**

Computation
 $s@t \rightarrow f(s) \rightarrow cmd@t'$
 $t' - t < bound$

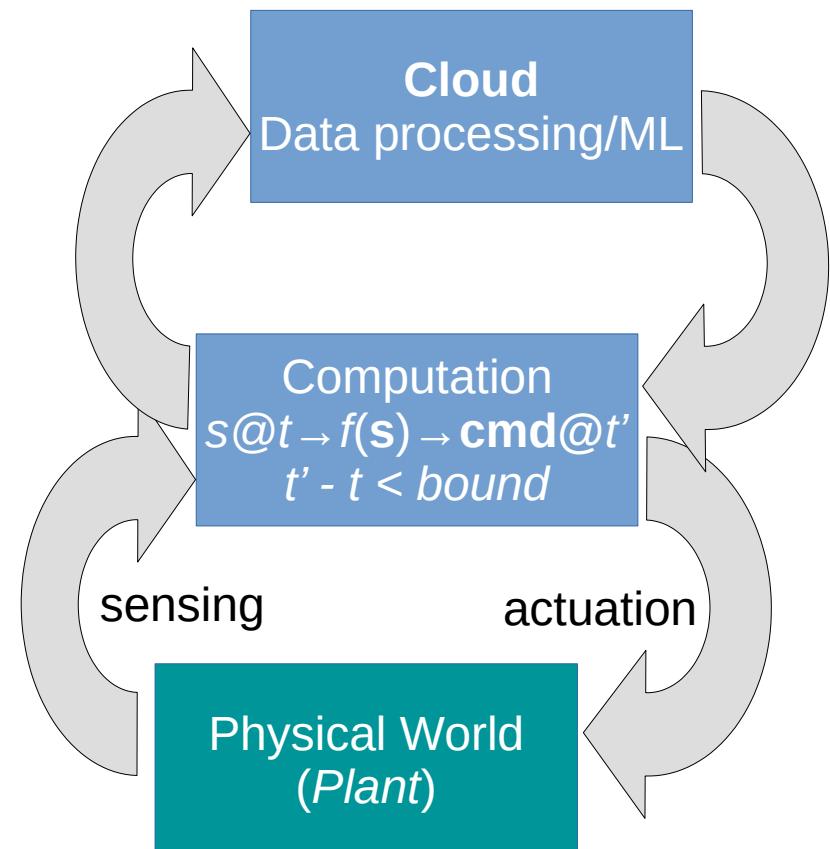
We don't write code thinking about time!

How do engineers program real-time systems???

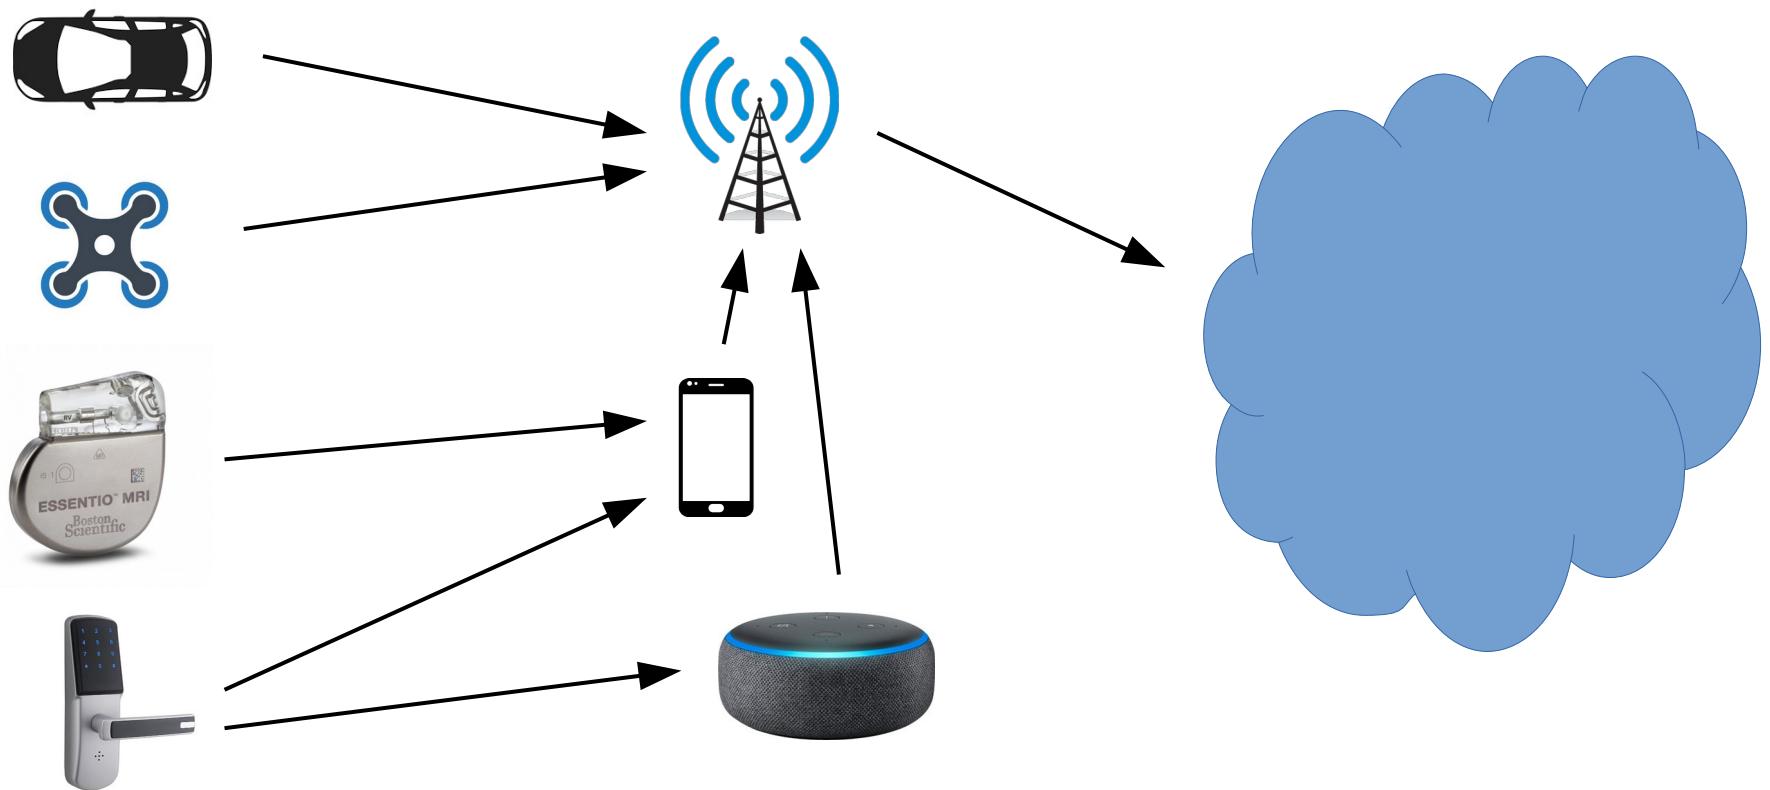
IoT Systems



Embedded/RT systems +
network



Vectors for IoT/CPS Attacks?



Building High-Confidence Embedded Systems

- How can we **design** for time?
- How can we **design** to handle malicious attacks/failures?

Perspective on Liability

Who is to blame if it goes wrong?

- Computer Science
 - Users = Lusers
 - “Move fast and **break things**”



Perspective on Liability

Who is to blame if it goes wrong?

- Computer Science
 - Users = Lusers
- Mechanical/civil/bio:
 - You build it,
you can be sued



Perspective on Liability

Who is to blame if it goes wrong?

- Computer Science
 - Autonomous Vehicles
 - Pervasive surveillance
 - City-level control



Now who's to blame? Regulation → Slower dev

Perspective on Liability

Who is to blame
when something goes wrong?

- Computer Science
 - Autonomous systems
 - Pervasive surveillance
 - City-level control

Goal: Can we *understand* the risk of the system, to explicitly *analyze* it?



Perspective on Liability

Who is to blame
when something goes wrong?

- Computer systems
 - Autonomous
 - Pervasive
 - City-level

Goal: Can we *understand* the risk
Really hard: Can we differentiate
between a *security breach*, a
programming error, and a
user error?



Perspective on Liability

Who is to blame
when something goes wrong

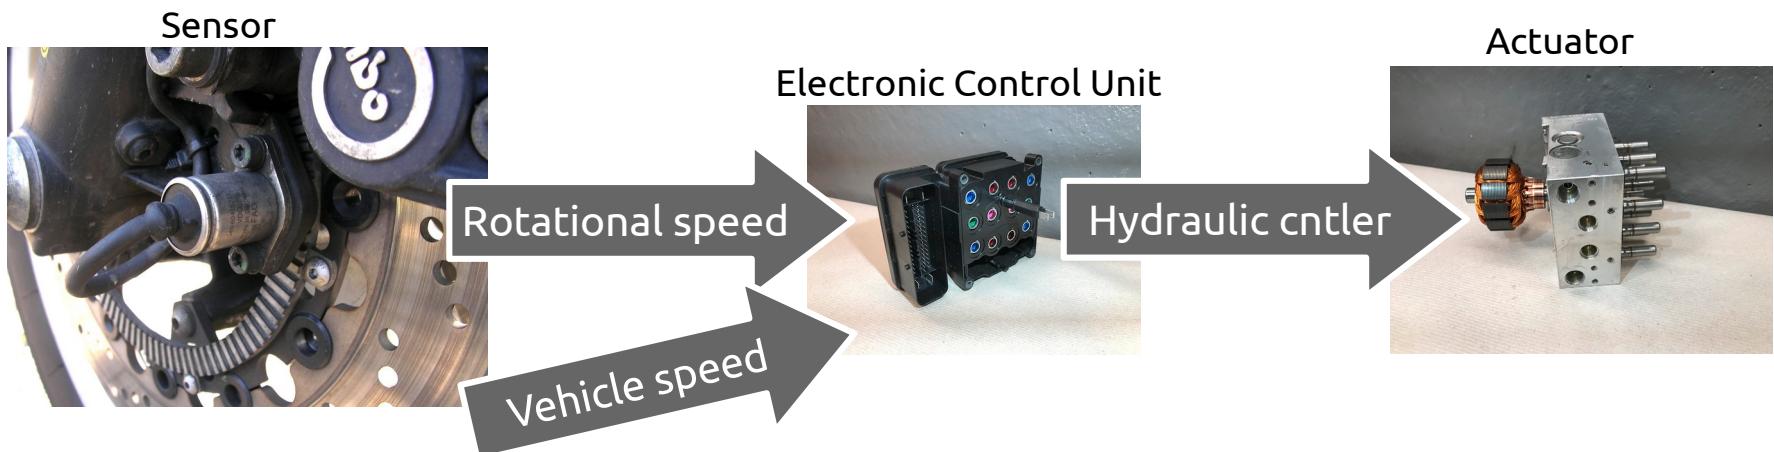
Goal: Can we *understand* the risk

- Compute **Really hard:** Can we differentiate
 - A
 - P
 - C
- **Really really hard:** Who is *legally culpable* for the damages (see Toyota uncontrolled acceleration case)



“Real-time” = time-awareness

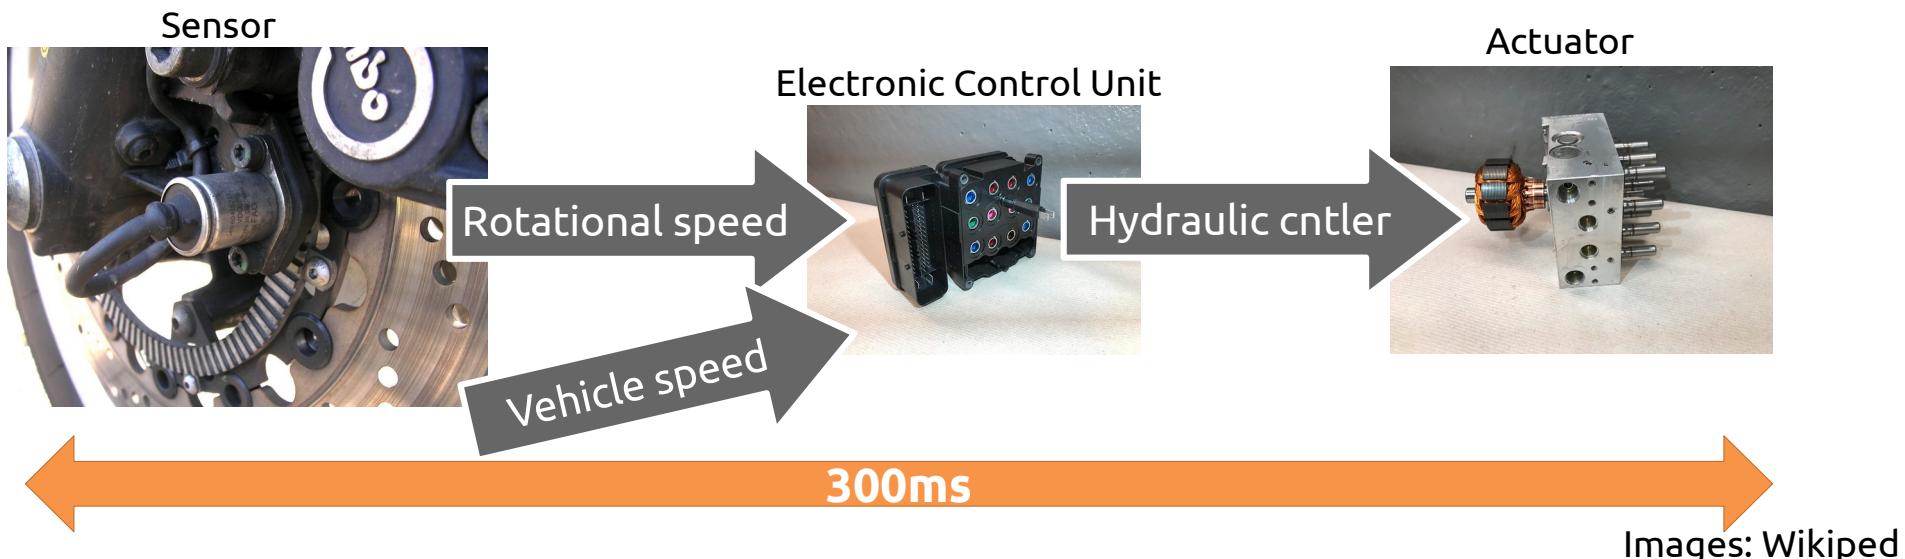
- Controlling physical world → Adhering to the world's timing constraints
- Example: Anti-lock brakes



Images: Wikipedia

“Real-time” = time-awareness

- Controlling physical world → Adhering to the **world's timing constraints**
- Example: Anti-lock brakes



“Real-time” = time-awareness

- Controlling physical world → Adhering to the **world's timing constraints**
- Example: Anti-lock brakes



RTSes: Deadline Examples

- Augmented Reality:
 - Human control → ~300ms (~150ms for pros)
 - human perception → ~10ms
- Robotics and AVs:
 - 1ms a robot needs to regulate force to the environment
 - <4ms walking robot applications
 - <10ms manipulating robots
 - <100ms driving robots
 - 10-100ms Vehicles: ATC, ABS

Real-Time Scheduling Goals

Goal: All tasks have **responsiveness** \leq deadline

Assumptions:

1. All threads execute **recurring code** (p).
2. We know apriori **how much computation** they do each activation (WCET).

Real-Time Scheduling Goals

```
while(1) {  
    // wakeup at next p  
    block_till_activation();  
    s = read_sensor();  
    o = computation(s);  
    update_actuator(o); // by deadline !  
}
```

deadline

1. All threads execute **recurring code** (p).
2. We know apriori **how much computation** they do each activation (WCET).

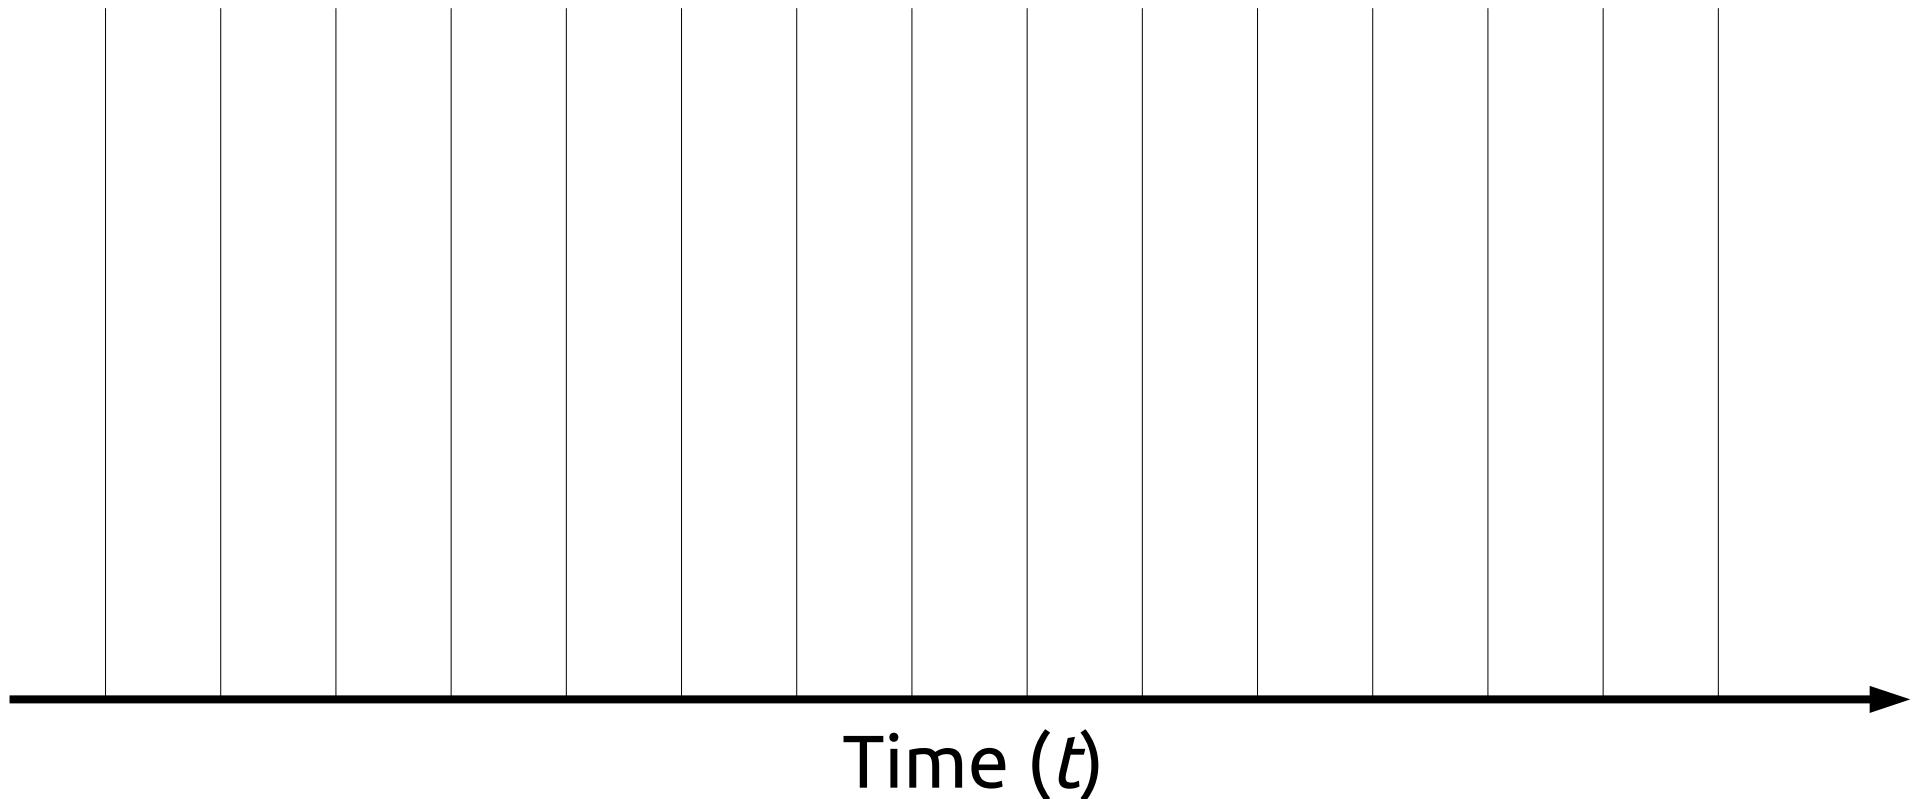
Periodic Task Model

Task (τ_i) – abstraction of computation, each has a

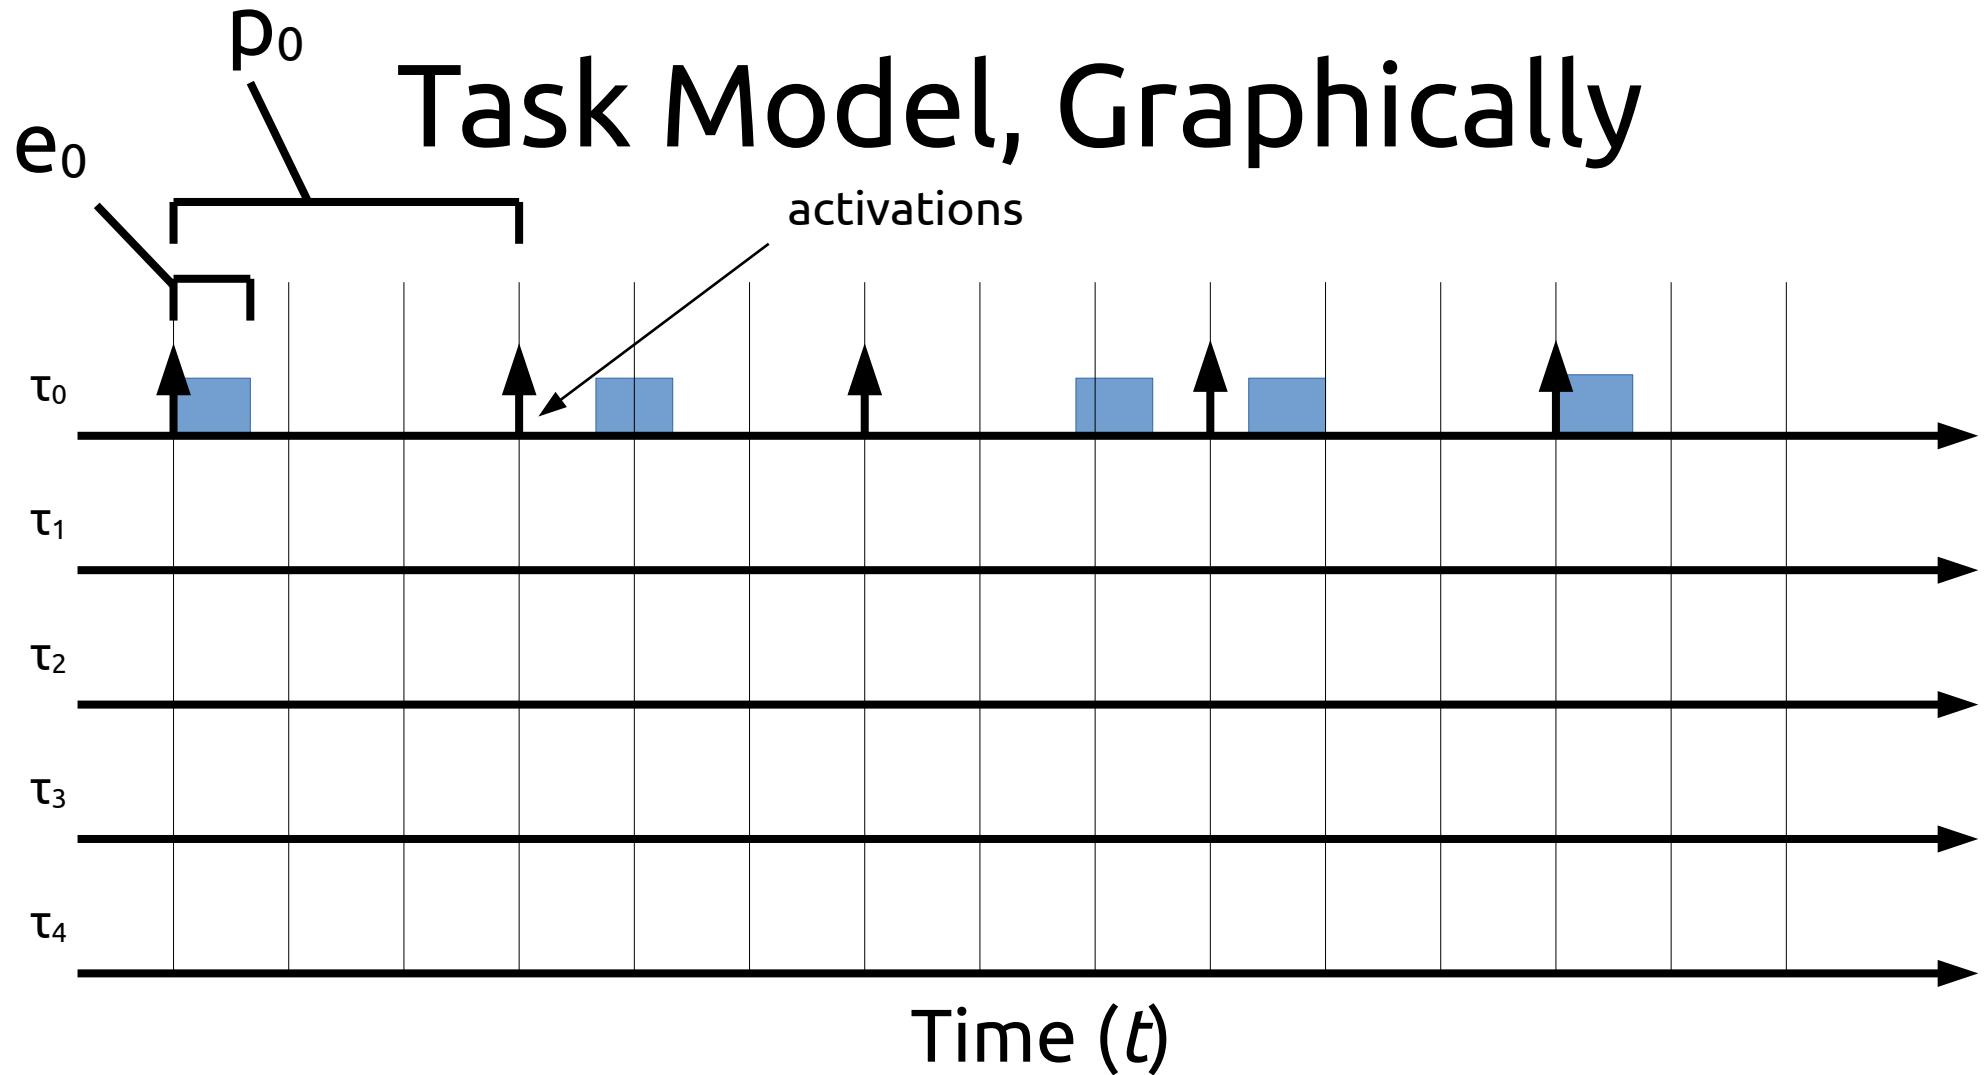
- **Period** (p_i) – period, span of repetitive exec.
A task *activates* (wakes up) every p_i time units
- **WCET** (e_i) – worst-case execution time
Every activation, task τ_i executes for max e_i
- **Deadline** of τ_i is p_i after its activation

System is $T = \{\tau_0, \tau_1, \dots, \tau_N\}$

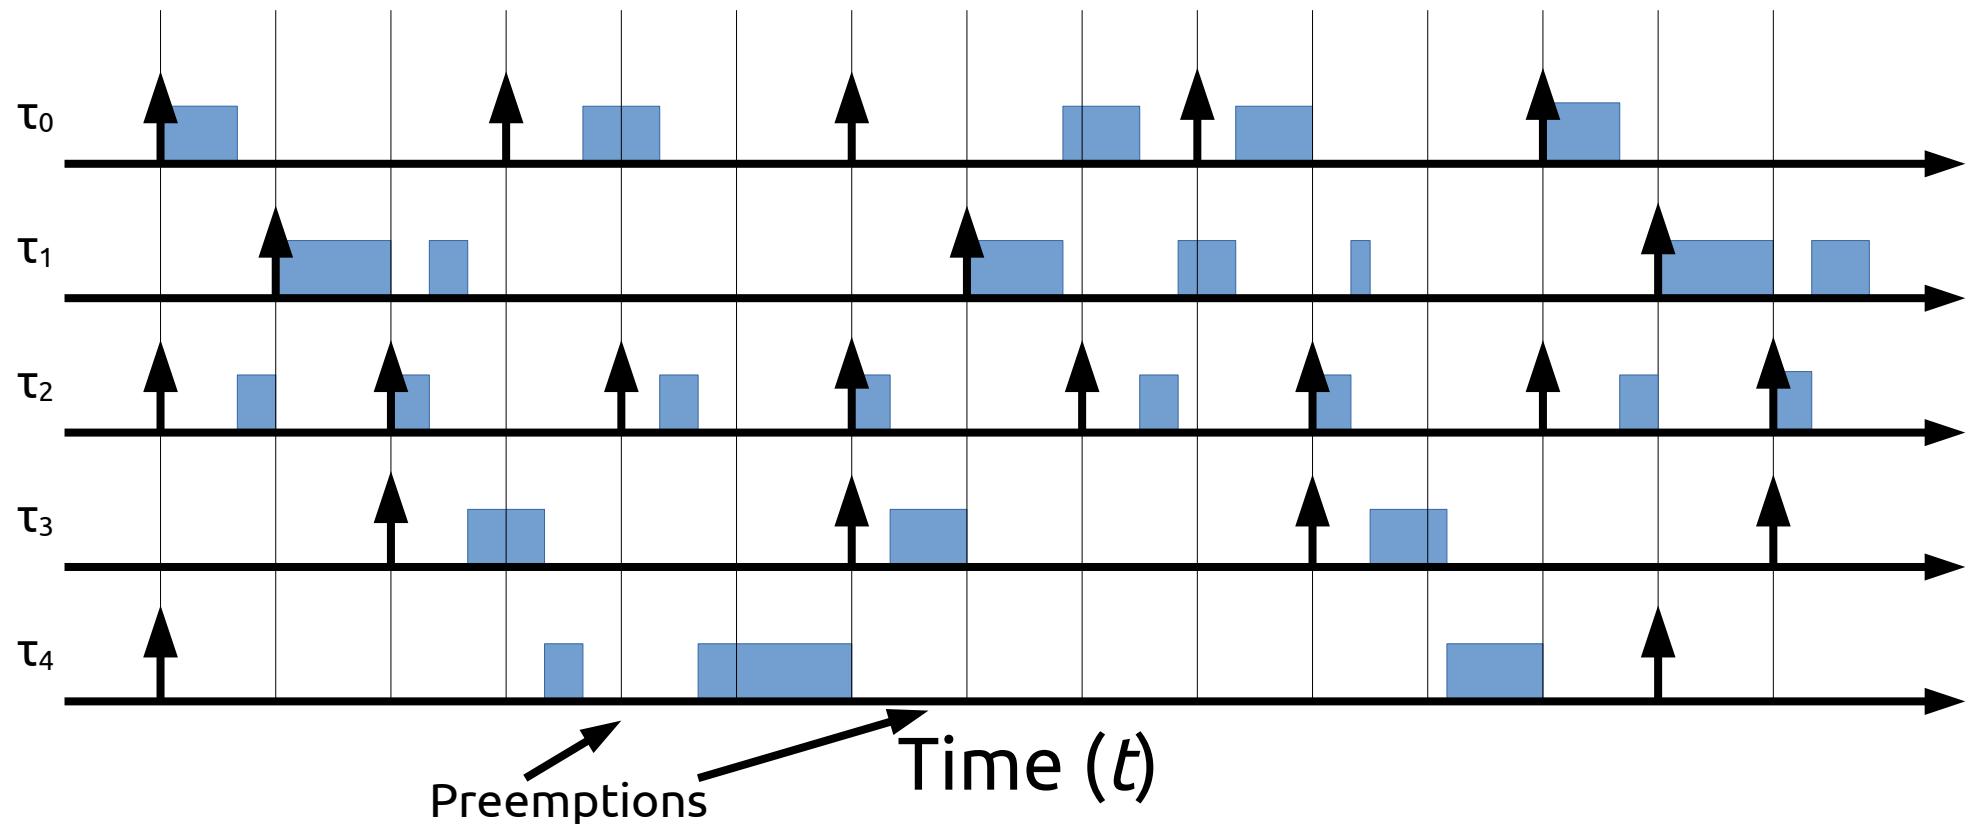
Task Model, Graphically

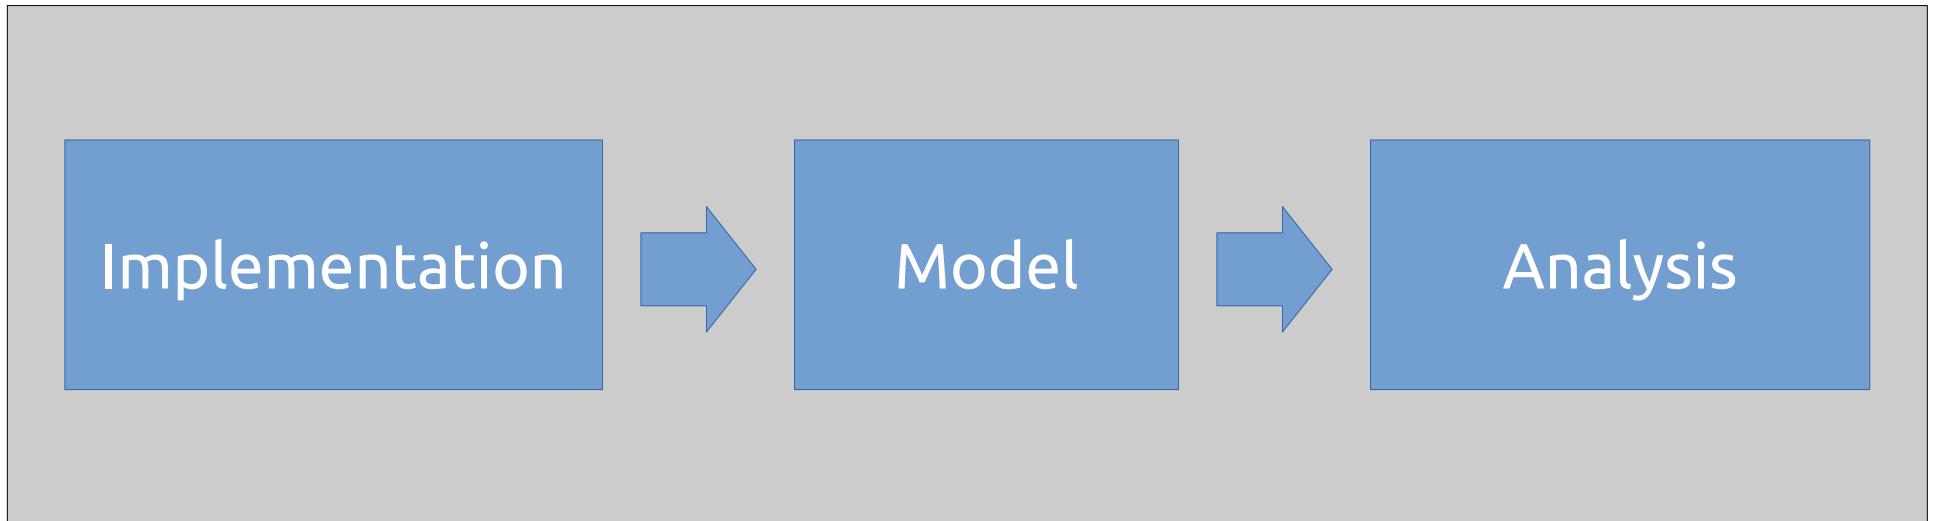


Task Model, Graphically



Task Model, Graphically





- *Goal:* real-time systems
 - **Implementation** → policies to instantiate a model
 - **Model** → mathematical description of impl.
 - **Analysis** → derive properties from model

Security/Confidence

- Network attached computers are vulnerable
 - No network → revert to '90s
 - **Plan for faults/compromises**

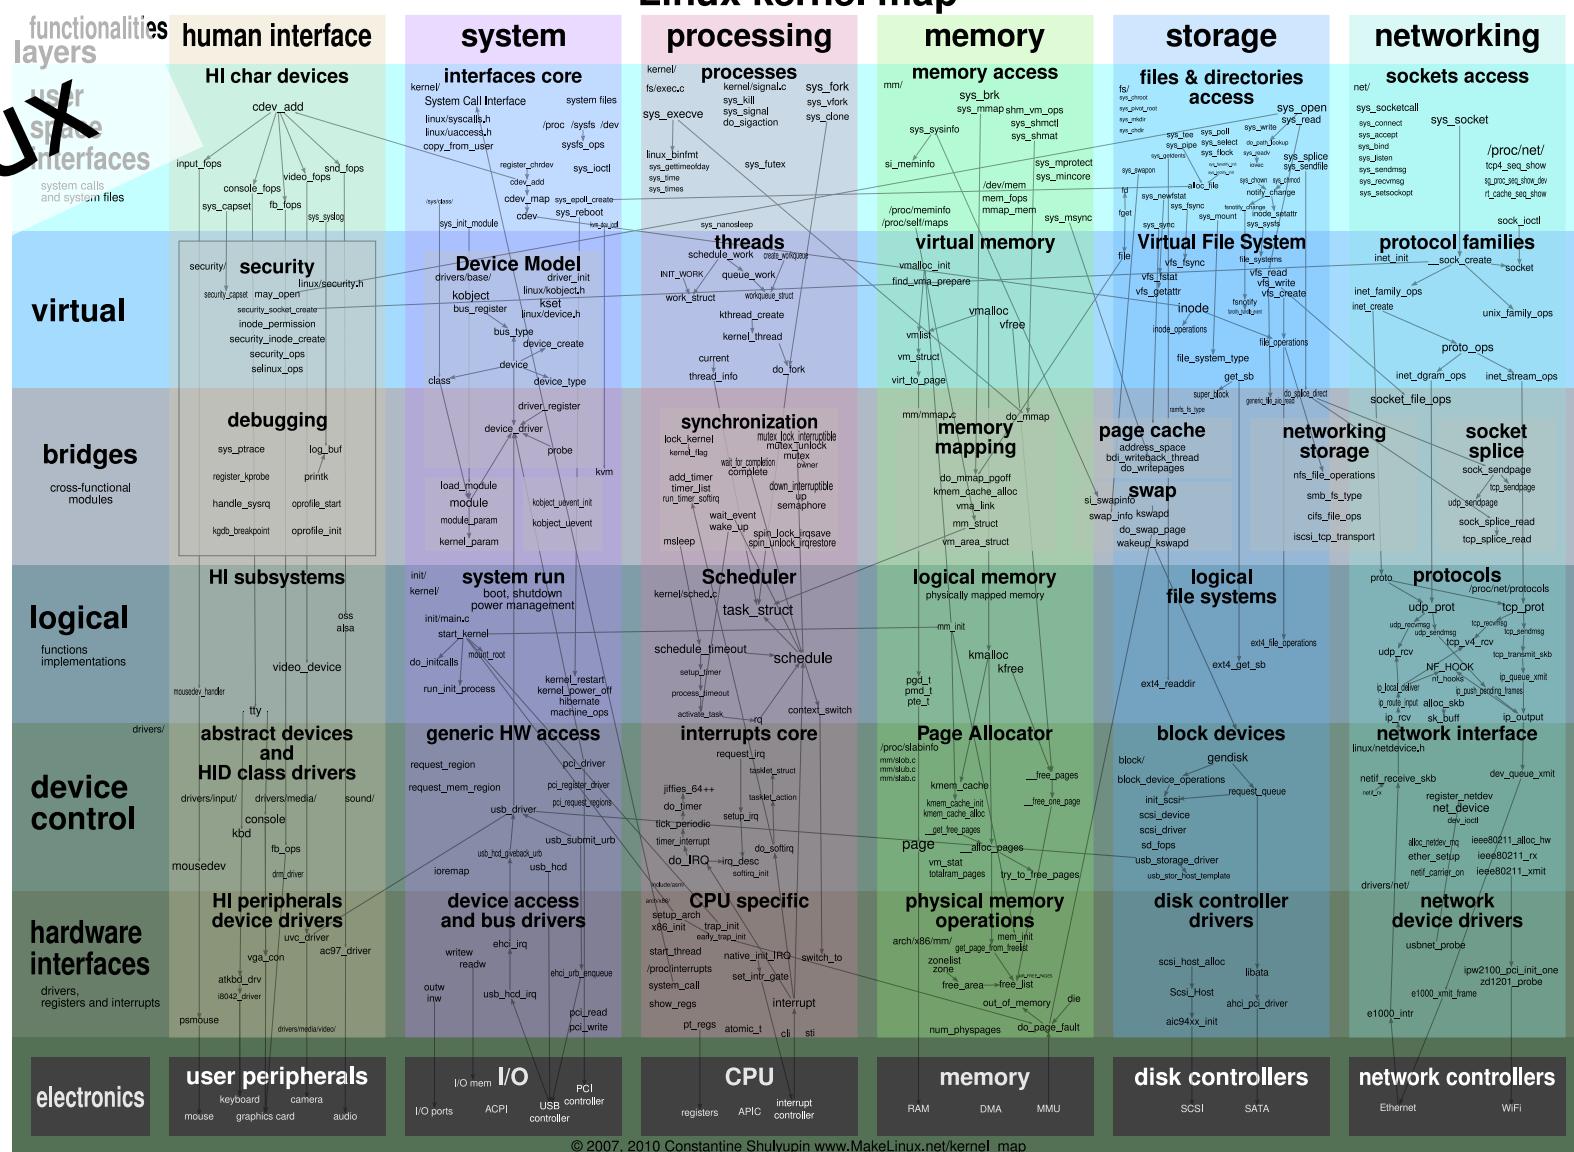
Security/Confidence

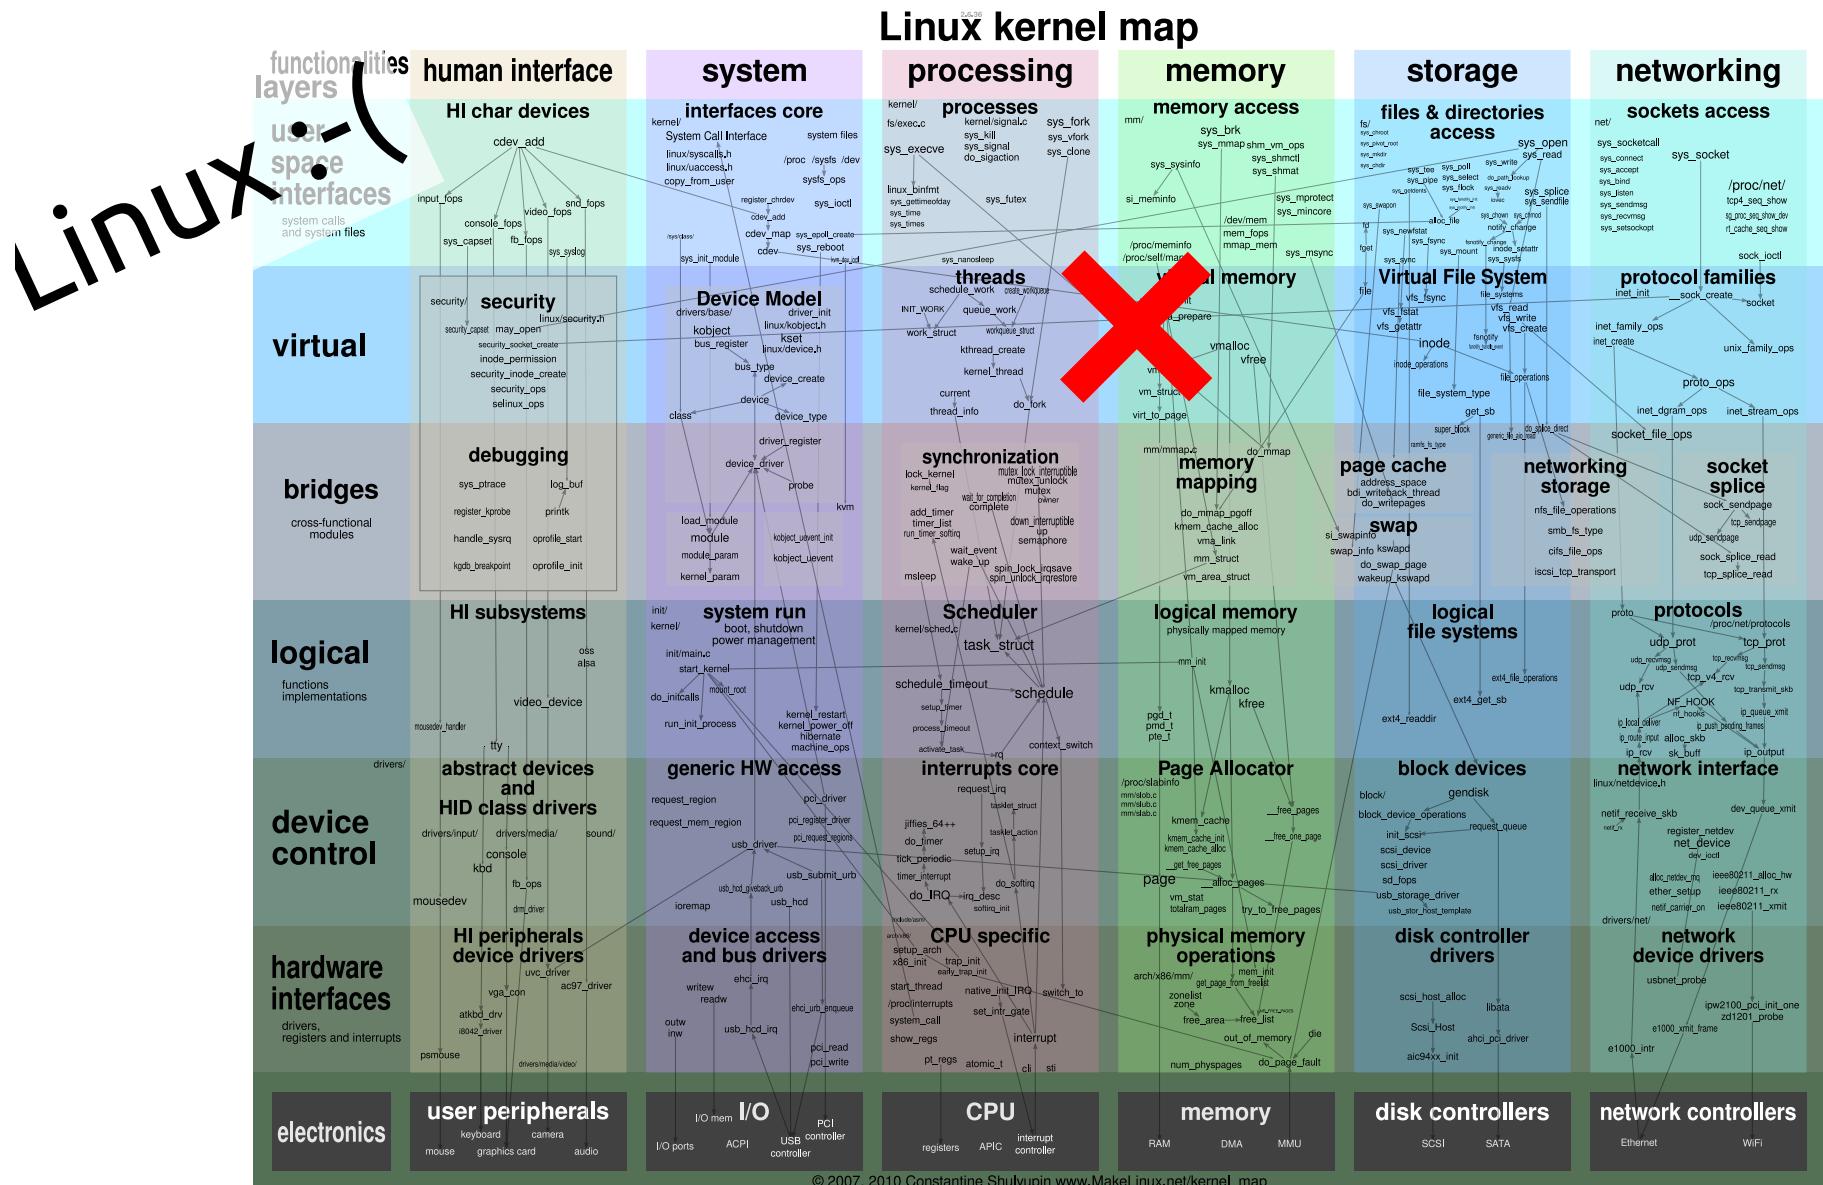
- Network attached computers are vulnerable
 - No network → revert to '90s
 - **Plan for faults/compromises**
- How?
 - **Isolation** → limit **scope** of effects of **faults**
 - **Controlled degradation** → faults lead to **lower functionality**, not **system failure**

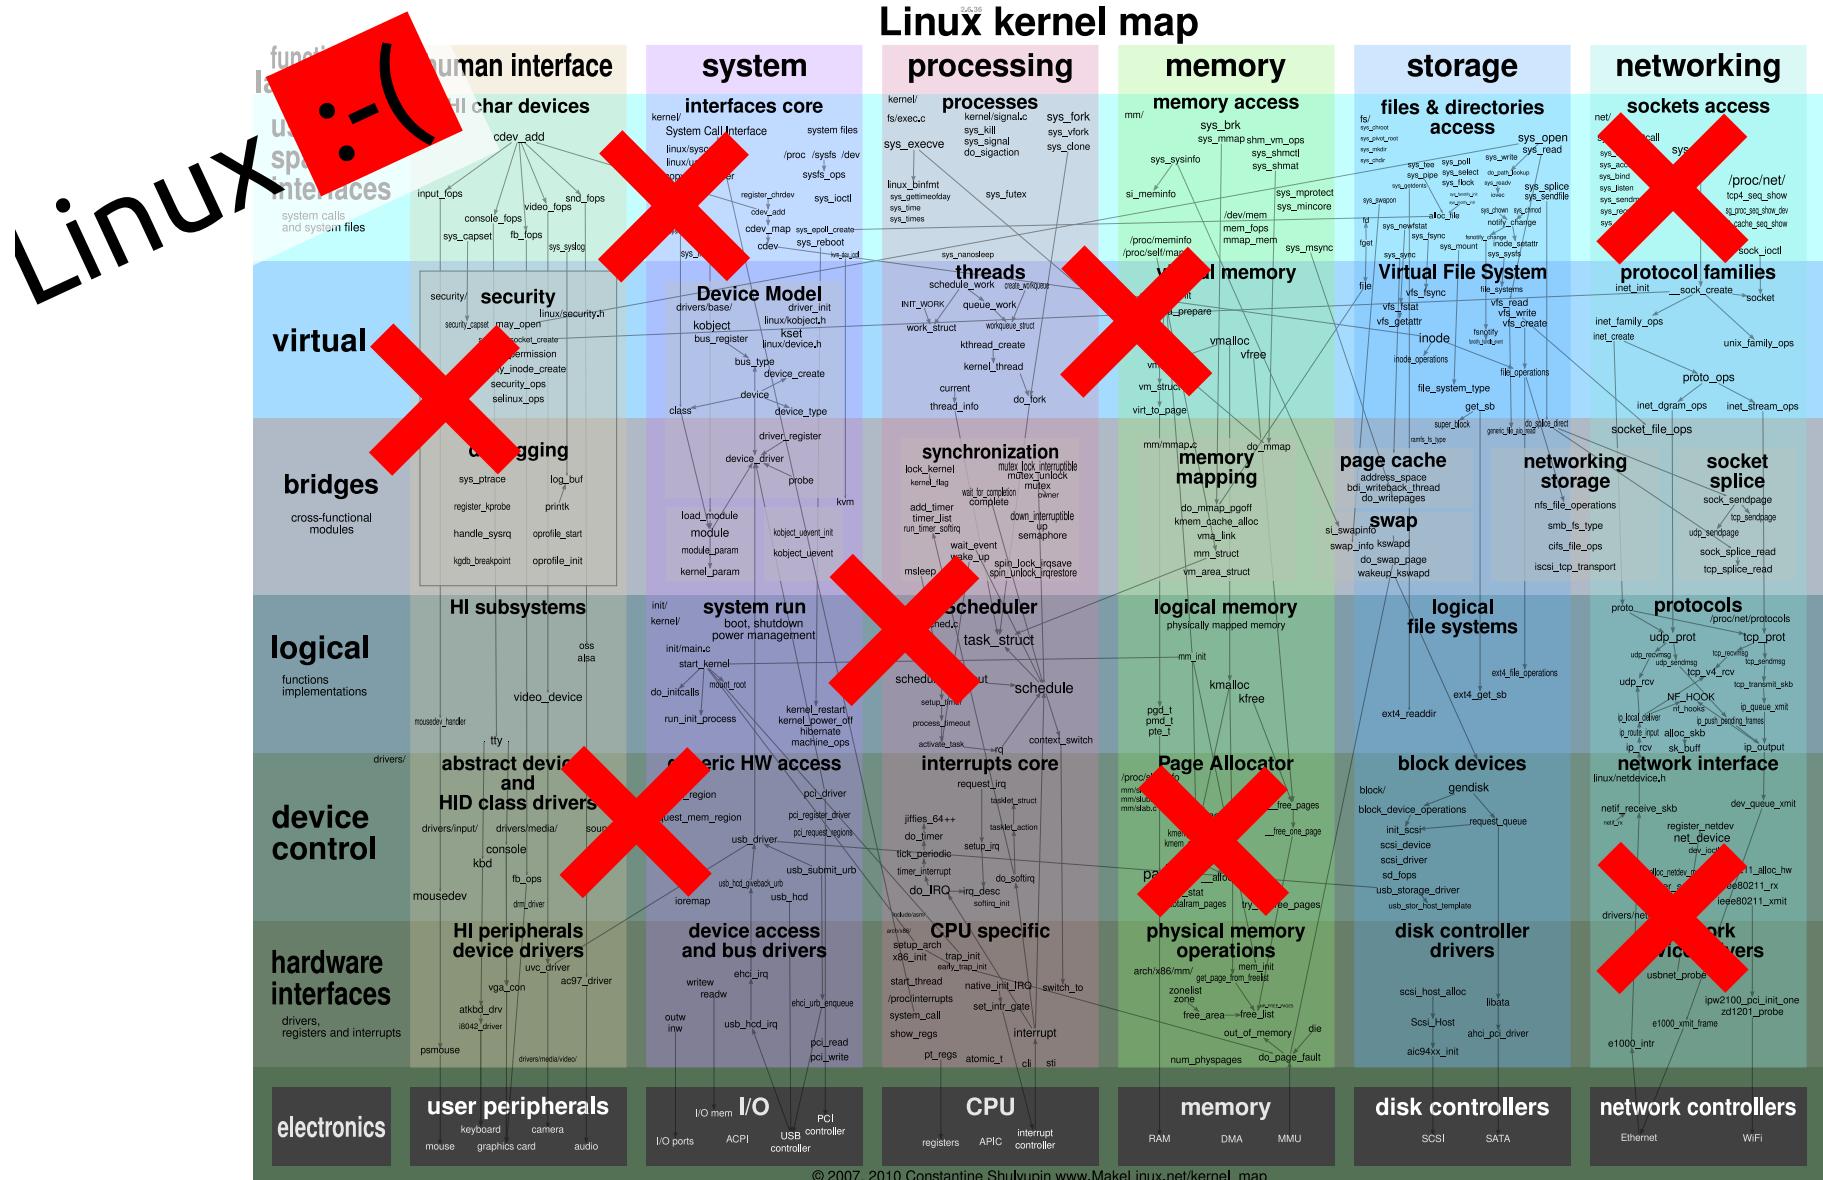
Operating Systems FTW

- Isolation in Linux
 - Processes can independently fail
 - No process can steal all CPU time

Linux





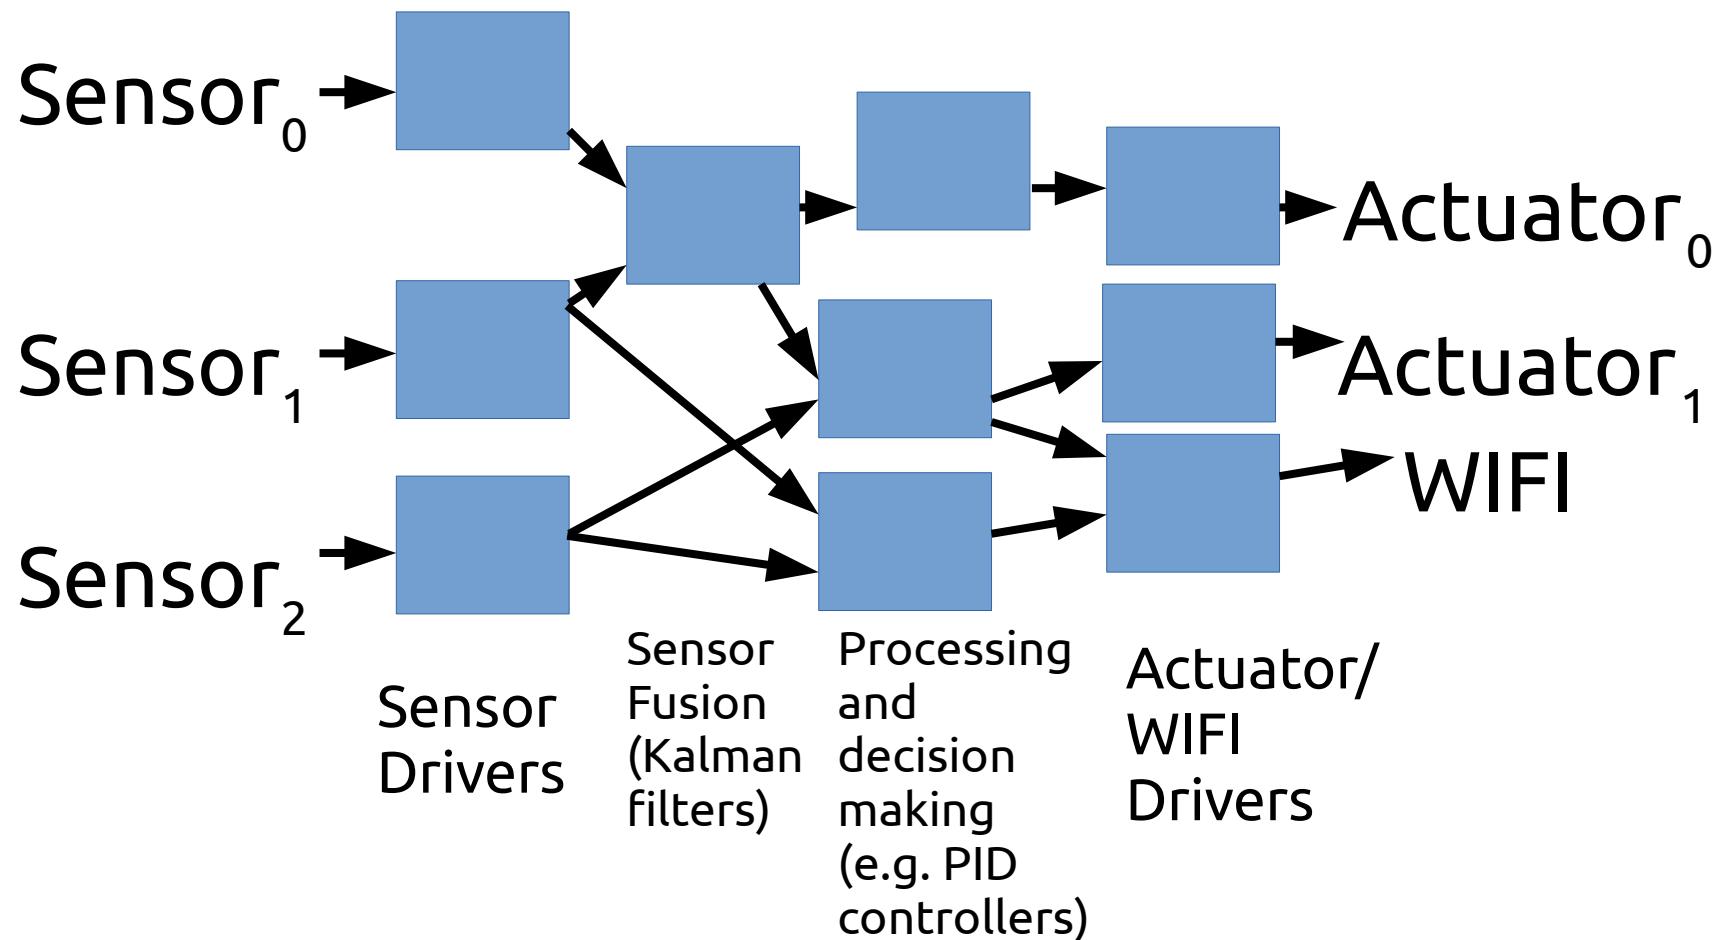


© 2007, 2010 Constantine Shulyupin www.MakeLinux.net/kernel_map

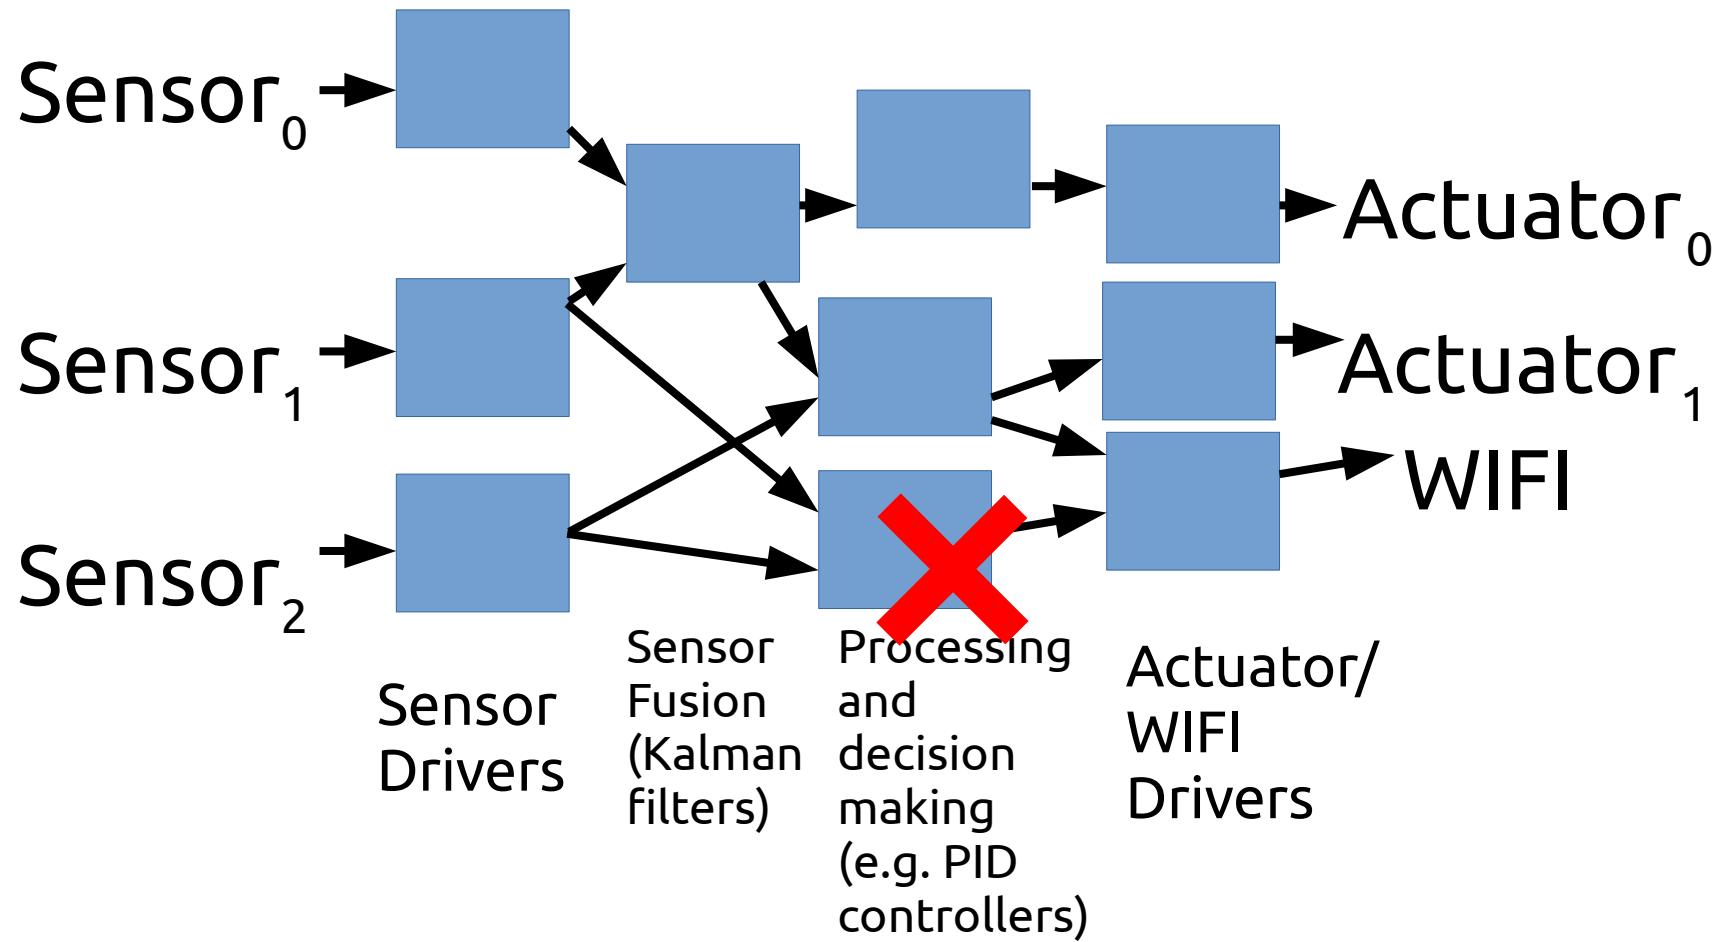
Secure Operating Systems

- **Break up** system software → minimize impact and scope of faults
- The “**core**” that provides isolation → must be **small/trustworthy**

Secure Operating Systems



Secure Operating Systems



High-Confidence Embedded Sys

- Real-time processing
- Isolation as a primary design concern